

# 3031\_A6\_Solutions\_v3

November 26, 2021

## 1 Assignment 6

(50 points total + 2 pts for naming/format)

### 1.1 Part 1

Read the spotify dataset from the file *spotify\_data.csv*.

What percentage of all the unique tracks are contributed by the top 3 artists of each genre, where the top artists are based on *artist\_popularity*, and the unique tracks are based on unique values of *track\_name*? (8 points for code)

A typical approach that will **not** work: If you group the data by genre, and filter the top 3 rows by *artist\_popularity*, then you may not get 3 unique artists, as one artist can have multiple tracks.

Here is one way to answer this question:

- (1) Group the data by genre, artist name and artist popularity. Find the number of unique tracks (by *track\_name*) for each group.
- (2) The dataset obtained in (1) is at artist-genre level, i.e., each row corresponds to a unique artist-genre combination. Group that dataset by genre, and filter the top 3 rows of each group based on artist popularity.
- (3) Sum up the number of unique tracks of the dataset obtained in (2) and divide it by the total number of unique tracks in the original dataset.

**Note:** (1) The functions *len()* and *unique()* will be useful.

(2) If you can propose a solution that is shorter than the one proposed above, on Monday - 15th Nov, in class, you will get 10% bonus points for this assignment.

```
[1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import warnings
#import seaborn as sns

warnings.filterwarnings('ignore')
#sns.set()
```

```
[2]: data = pd.read_csv('spotify_data.csv')
```

```
[3]: grouped = data.groupby(['genres', 'artist_name', 'artist_popularity'])
```

```
[4]: grouped['track_name'].nunique().reset_index()
```

```
[4]:
```

	genres	artist_name	artist_popularity	track_name
0	country	Alan Jackson	76	61
1	country	Blake Shelton	77	16
2	country	Brad Paisley	71	16
3	country	Brett Young	73	5
4	country	Casey Donahew	60	3
...	...	...	...	...
11949	rock	Óscar Chávez	51	33
11950	rock	İkiye On Kala	61	8
11951	rock		52	52
11952	rock	(!)	58	22
11953	rock		57	1

```
[11954 rows x 4 columns]
```

```
[5]: artist_genre_lvl=grouped['track_name'].nunique().reset_index()
top3_artists_data = artist_genre_lvl.groupby('genres').apply(lambda x:x.
    ↳sort_values(by = ['artist_popularity'],ascending = False)[0:3])
top3_artists_data.track_name.sum()/len(data.track_name.unique())
```

```
[5]: 0.047140401953927644
```

The top 3 artisits of each genre contribute to 5% of the total number of tracks.

## 1.2 Part 2

Read data from the file “Canadian\_Fish\_Biodiversity.csv” on Canvas. Each row records a unique fishing event from a 2013 sample of fish populations in Ontario, Canada. *(42 points overall)*

```
[6]: cfldata = pd.read_csv("Canadian_Fish_Biodiversity.csv")

cfldata["Species"].nunique()
```

```
[6]: 132
```

### 1.2.1 Question 1

To analyze the results of these fishing surveys, we need to understand the dynamics of projects, sites, and geographic locations. In large part the following questions deal with missing data. *(16 points total)*

- Each site (identified by the column *SITEID*) represents a time and place at which fishing events occurred. Sites are grouped into broader projects (identified by the column *Project Name*). We want to understand the scope of these projects.

Using *.groupby*, find the top three projects by number of unique sites. *(2 points for code)*

**Hint:** The Pandas function `nunique()` may help

```
[7]: cfbddata.groupby("Project Name")["SITEID"].nunique().  
      ↪sort_values(ascending=False).head(3)
```

```
[7]: Project Name  
      2013 GLAP Survey of Detroit River      220  
      2013 Crown Marsh Survey                146  
      2013 Spotted Gar Critical Habitat Survey 131  
      Name: SITEID, dtype: int64
```

b) Find the top three and bottom three projects in terms of the proportion of unique sites of the total number of unique sites. (3 points for code)

```
[8]: #top 3  
(cfbddata.groupby("Project Name")["SITEID"].nunique()/\  
cfbddata.groupby("Project Name")["SITEID"].count()).sort_values().head(3)
```

```
[8]: Project Name  
      2013 Grass Pickerel Twenty Mile Creek      0.047619  
      2013 Mussel Fish Community Assessment      0.056452  
      2013 Lake Chubsucker Critical Habitat Survey 0.056572  
      Name: SITEID, dtype: float64
```

```
[9]: #bottom 3  
(cfbddata.groupby("Project Name")["SITEID"].nunique()/\  
cfbddata.groupby("Project Name")["SITEID"].count()).sort_values().tail(3)
```

```
[9]: Project Name  
      2013 Eastern Sand Darter eDNA Survey of Sydenham River      1.0  
      2013 Eastern Sand Darter eDNA Survey of Grand River        1.0  
      2013 Spotted Gar eDNA Survey                               1.0  
      Name: SITEID, dtype: float64
```

c) (i) How many values are missing for the air temperature column? (1 point for code)

```
[10]: cfbddata['Air Temperature (C)'].isnull().sum()
```

```
[10]: 808
```

(i) 808 values are missing for the air temperature column.

(ii) Impute the missing values of air temperature with the median air temperature of the corresponding water body (*Waterbody Name*) and month. (2 points for code)

(ii)

```
[11]: cfbddata["Air Temperature (C)"] = cfbddata.groupby(["Waterbody_\  
      ↪Name", 'Month'])["Air Temperature (C)"].apply(lambda x:x.fillna(x.median()))
```

- (iii) How many missing values still remain for the air temperature column after the imputation in (ii)? (1 point for answer)

```
[12]: cfbddata["Air Temperature (C)"].isnull().sum()
```

```
[12]: 113
```

- (iii) 113 missing values still remain after the imputation in (ii)

- (iv) We will try to impute the remaining missing values for air temperature. Try impute the remaining missing values of air temperature with the median air temperature of the corresponding project (*Project Name*) and month. (2 points for code)

(iv)

```
[13]: cfbddata["Air Temperature (C)"] = cfbddata.groupby(["Project Name", 'Month'])["Air_
      ↪Temperature (C)"].apply(lambda x:x.fillna(x.median()))
```

- (v) How many missing values still remain for the air temperature column after the imputation in (iv)? (1 point for answer)

```
[14]: cfbddata["Air Temperature (C)"].isnull().sum()
```

```
[14]: 62
```

- (v) 62 missing values still remain after the imputation in (iv)

- (vi) Find the correlation between air temperature and water temperature. (1 point for code)

(vi)

```
[15]: cfbddata["Air Temperature (C)"].corr(cfbddata["Water Temperature (C)"])
```

```
[15]: 0.768184572633517
```

Correlation = 77%

- (vii) As you found a high correlation between air temperature and water temperature in (vi), you can use water temperature to estimate the air temperature (using the trendline, like you did in assignment 4). Assuming you already did that, how many missing values will still remain for the air temperature column? *Note: Do not impute the missing values using the trendline, just assume you already did that.* (1 point for code)

- (vii) The values for air temperature will remain missing for those observations that have missing values of water temperature.

```
[16]: (cfbddata['Air Temperature (C)'].isnull() & cfbddata['Water Temperature (C)'].
      ↪isnull()).sum()
```

```
[16]: 11
```


11 missing values will still remain after imputing the missing values using the trendline.

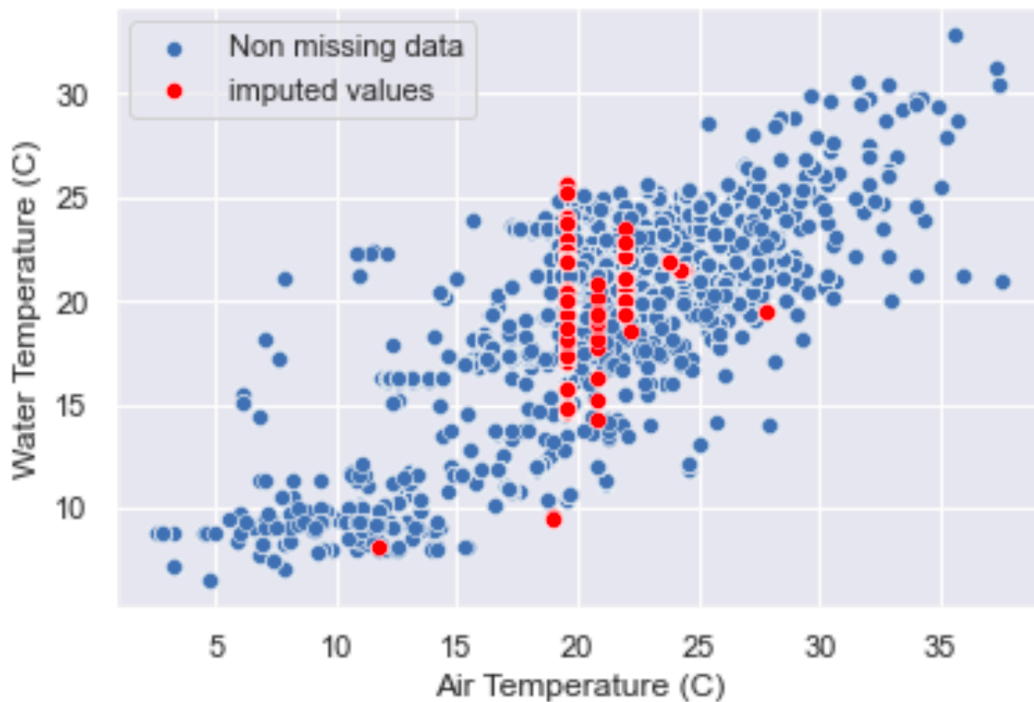
(viii) Make a scatterplot of air temperature against water temperature. Highlight the points for which the air temperature was imputed in (ii) and (iv) with a different color. (2 points for code)

(viii)

```
[17]: cfbddata_orig = pd.read_csv("Canadian_Fish_Biodiversity.csv")

[18]: ind_imputed = np.where(cfbddata_orig['Air Temperature (C)'].isnull() &
    ~cfbddata['Air Temperature (C)'].isnull())[0]
plt.scatter(data = cfbddata_orig, x = 'Air Temperature (C)', y = 'Water_
    ~Temperature (C)')
plt.scatter(data = cfbddata.iloc[ind_imputed,:], x = 'Air Temperature (C)', y =
    ~'Water Temperature (C)',color = 'red')
plt.legend(labels = ['Non missing data','imputed values'])
plt.show()
```





### 1.2.2 Question 2

This section begins to investigate the living conditions of fish at different locations and time periods. (7 points total)

- a) Use a single `.groupby` statement to view the minimum, mean, standard deviation, and maximum air temperature and water temperature for each project during the month of August (use the `Month` column). (2 points for code)

```
[19]: cfb_month = cfbddata[cfbddata["Month"]==8]
cfb_month.groupby("Project Name")[["Air Temperature (C)", "Water Temperature (C)"]].agg(['min', 'mean', 'std', 'max'])
```

```
[19]:
```

	Air Temperature (C)		
	min	mean	\
Project Name			
2013 Bridle Shiner Critical Habitat Survey	20.7	24.609091	
2013 Crown Marsh Survey	16.4	21.673275	
2013 GLAP Survey of Detroit River	21.1	24.360619	
2013 Grass Pickerel Niagara Drains	20.8	25.246154	
2013 Grass Pickerel Twenty Mile Creek	22.5	25.226190	
2013 Lake Chubsucker Critical Habitat Survey	14.3	21.136106	
2013 Mussel Fish Community Assessment	23.0	24.535887	
2013 Pugnose Minnow Lake St Clair Drains	22.2	25.303061	
2013 Species at Risk Assessment	23.3	24.793939	
2013 Spotted Gar Critical Habitat Survey	18.6	22.706481	

	std	max
Project Name		
2013 Bridle Shiner Critical Habitat Survey	2.588098	26.5
2013 Crown Marsh Survey	2.040997	26.2
2013 GLAP Survey of Detroit River	2.555656	28.3
2013 Grass Pickerel Niagara Drains	2.498615	28.2
2013 Grass Pickerel Twenty Mile Creek	2.283649	29.4
2013 Lake Chubsucker Critical Habitat Survey	3.427318	32.6
2013 Mussel Fish Community Assessment	1.203286	26.1
2013 Pugnose Minnow Lake St Clair Drains	1.961982	29.0
2013 Species at Risk Assessment	0.559087	25.0
2013 Spotted Gar Critical Habitat Survey	2.590503	27.7

	Water Temperature (C)	
	min	mean
Project Name		
2013 Bridle Shiner Critical Habitat Survey	20.630	22.052182
2013 Crown Marsh Survey	17.430	22.007084
2013 GLAP Survey of Detroit River	21.646	22.028226
2013 Grass Pickerel Niagara Drains	17.150	20.105641
2013 Grass Pickerel Twenty Mile Creek	18.990	22.380238
2013 Lake Chubsucker Critical Habitat Survey	19.300	21.671746
2013 Mussel Fish Community Assessment	21.800	23.155645
2013 Pugnose Minnow Lake St Clair Drains	18.970	20.713163
2013 Species at Risk Assessment	21.960	22.075600
2013 Spotted Gar Critical Habitat Survey	18.639	21.349769

	std	max
Project Name		
2013 Bridle Shiner Critical Habitat Survey	0.503887	22.44
2013 Crown Marsh Survey	1.240151	23.24
2013 GLAP Survey of Detroit River	0.392402	23.11
2013 Grass Pickerel Niagara Drains	2.426792	23.96
2013 Grass Pickerel Twenty Mile Creek	2.533686	26.87
2013 Lake Chubsucker Critical Habitat Survey	1.619878	25.13
2013 Mussel Fish Community Assessment	1.156345	24.60
2013 Pugnose Minnow Lake St Clair Drains	1.676571	24.22
2013 Species at Risk Assessment	0.080936	22.13
2013 Spotted Gar Critical Habitat Survey	1.978877	24.30

b) Make lineplots showing maximum air temperature and water temperature by month and *Region*. To construct *Region*, use *pd.cut* to satisfy the following conditions:

- Rows with a latitude lower than 42.4 should have *Southern* in the *Region* column
- Rows with a latitude between 42.4 and 42.8 should have *Central* in the *Region* column

- Rows with a latitude higher than 42.8 should have *Northern* in the *Region* column


You can have the month on the horizontal axis, the temperature on the vertical axis, different colors for different regions, and different styles (solid line / dotted line) to indicate air/water temperature.

Does anything in the visualization surprise you? Why or why not? (4 points for code and visualization, 1 point for answer)

```
[20]: lat_bins = [40,42.4,42.8,46]
lat_names = ["Southern", "Central", "Northern"]

cfbdata["Region"] = pd.cut(cfbdata["Start Latitude"], lat_bins, labels=lat_names)
```

```
[21]: datag = cfbdata.groupby(['Region', 'Month'], as_index = False)[['Air Temperature_
    ↳(C)', 'Water Temperature (C)']].max()
data_melt = pd.melt(datag, id_vars = ['Region', 'Month'], var_name = 'Temperature_
    ↳type', value_name = 'Temperature measurement')
#sns.boxplot(data = data_melt, x = 'Month', y = 'Temperature measurement', hue =
    ↳'Region', style = 'Temperature type')
#plt.legend(bbox_to_anchor = (1,1))
```



(Sample answer – any reasonable interpretation is acceptable)

I'm surprised that both water and air maximum temperatures in the Northern region are higher than other regions' in June and July—I would've expected bodies of water further north to be generally colder. I'm also surprised that maximum temperatures in the Central region increase from September to October.

### 1.2.3 Question 3

Finally let's focus on the stars of this survey—the fish, of course. (19 points total)



- a) Let's continue using our *Region* categorization. Find the top three fish species in each region by number captured. (3 points for code)

```
[22]: top3caught = lambda c: c.sort_values(by="Number Captured",ascending=False).
      ↪head(3)

      (cfbdata.groupby(["Region","Species"],as_index=False)["Number Captured"]\
      .sum()).groupby("Region", as_index=False).apply(top3caught)
```

```
[22]:
```

		Region	Species	Number Captured
0	63	Southern	Lepomis macrochirus	5072.0
	26	Southern	Dorosoma cepedianum	3559.0
	82	Southern	Neogobius melanostomus	2265.0
1	195	Central	Lepomis macrochirus	2126.0
	186	Central	Labidesthes sicculus	2029.0
	223	Central	Notropis heterodon	1562.0
2	346	Northern	Neogobius melanostomus	2522.0
	362	Northern	Notropis volucellus	2104.0
	327	Northern	Lepomis macrochirus	1841.0



- b) Are certain fish only found in some regions? Visualize how many species are in all three regions, how many are in two of three, and how many were only captured in one region. (3 points for code and visualization)

```
[23]: region_num = pd.DataFrame(cfbdata.groupby("Species",as_index=False)["Region"]
      ↪.nunique()["Region"].value_counts()).reset_index()
      region_num.columns = ["Number of Regions","Number of Species"]
      region_num
```

```
[23]:
```

	Number of Regions	Number of Species
0	3	51
1	1	51
2	2	30

```
[24]: plt.bar(x="Number of Regions",height="Number of Species",data=region_num)
      plt.show()
```



- c) What percentage of all species are exclusively captured in the Southern region? How about the Northern Region? And the Central region? (3 points for code)

```
[25]: #Method 1
exclusive_species = pd.DataFrame(cfbddata.
    ↳groupby("Species",as_index=False)["Region"].nunique())
exclusive_species = exclusive_species[exclusive_species["Region"]==1]

exclusive_species.head()
```

```
[25]:
```

	Species	Region
5	Ameiurus sp	1
10	Carassius auratus X Cyprinus carpio	1
12	Catostomidae	1
16	Chrosomus eos	1
17	Clupeidae	1

```
[26]: cfbddata["Species"].nunique()
```

```
[26]: 132
```

```
[27]: exclusive_obs = cfbddata[cfbddata["Species"].isin(exclusive_species["Species"])]
exclusive_obs.groupby("Region")["Species"].nunique() #/cfbddata["Species"].
    ↳nunique()
```

```
[27]: Region
Southern    12
Central      6
Northern    33
Name: Species, dtype: int64
```

```
[28]: #Method 2:
dp = cfbddata.pivot_table(index = 'Species', columns = 'Region', values = _
    ↳'Number Captured',aggfunc = 'sum',margins = True)
for reg in ['Northern','Southern','Central']:
    print('Region: ',reg,":",((dp[reg] == dp['All']) & (dp[reg]>0)).sum()/
    ↳cfbddata.Species.nunique())
```

```
Region: Northern : 0.25
```

```
Region: Southern : 0.09090909090909091
```

```
Region: Central : 0.045454545454545456
```

d) Turbidity quantifies the level of cloudiness in liquid. For fish in each of the three regions, is there a correlative relationship between turbidity and # of fish caught? (2 points for code, 1 point for answer)

```
[29]: cfbddata.groupby('Region').apply(lambda x:x['Turbidity (ntu)'].corr(x['Number_
    ↳Captured']))
```

```
[29]: Region
Southern    -0.019202
```

```
Central      -0.016327
Northern     0.063456
dtype: float64
```

No, there does not appear to be a correlation.

- e) Now let's turn to the length of fish captured, given by *Maximum (mm)* and *Minimum (mm)*. Find the overall maximum and minimum lengths of all fish in each region. Which region has the largest range in captured fish length? (2 points for code, 1 point for answer)

```
[30]: cfbdata.groupby("Region").agg({"Maximum (mm)": 'max', "Minimum (mm)": 'min'})
```

```
[30]:      Maximum (mm)  Minimum (mm)
Region
Southern          1130.0           8.0
Central            785.0           9.0
Northern           760.0          10.0
```

The Southern region has the largest range in captured fish length.

- f) Find the inverse Simpson index of species counts for each waterbody type (*WaterbodyType*) within each region. Which combination of waterbody type and region has the greatest diversity of fish species? Which has the least?

The inverse Simpson index ( $\frac{1}{\lambda}$ ) is a measure of ecological diversity, for which a larger index number indicates a greater diversity of species. The index is calculated as:

$$\frac{1}{\lambda} = 1 / (\sum_{i=1}^R p_i^2)$$

where  $R$  is the number of unique species and  $p_i$  is the proportion of fish belonging to species  $i$ . (3 points for code, 1 point for answer)

```
[31]: #method 1:
wrs_captured = cfbdata.
    ↳groupby(["WaterbodyType", "Region", "Species"], as_index=False) ["Number_
    ↳Captured"].sum()
wrs_captured = wrs_captured[wrs_captured["Number Captured"]>0]

inv_simpson = lambda s: 1/(np.sum((s/np.sum(s))**2))

wrs_captured.groupby(["WaterbodyType", "Region"])["Number Captured"].
    ↳apply(inv_simpson)
```

```
[31]: WaterbodyType  Region
Lake              Southern    4.094649
              Central      3.854458
              Northern    4.896119
Stream           Southern   13.800772
              Central    12.070477
              Northern   20.306997
Wetland          Southern      NaN
```

```

                Central      8.954213
                Northern      NaN
Name: Number Captured, dtype: float64

```

```

[32]: #Method 2:
def f(x):
    R = x.Species.nunique()
    all_num = x['Number Captured'].sum()
    all_species = x.Species.unique()
    si=0
    for i in all_species:
        icount = x.loc[x['Species']==i, 'Number Captured'].sum()
        if icount>0:
            si = si+((icount/all_num)**2)
    return 1/si

```

```

[33]: cfbdata.groupby(['WaterbodyType', 'Region']).apply(f)

```

```

[33]: WaterbodyType  Region
Lake              Southern    4.094649
              Central      3.854458
              Northern    4.896119
Stream           Southern   13.800772
              Central    12.070477
              Northern   20.306997
Wetland          Central     8.954213
dtype: float64

```

According to the inverse Simpson index, Northern streams have the greatest diversity in fish species and Central lakes have the least. The survey didn't include any fish in Southern or Northern wetlands, so we can't describe these bodies of water in terms of fish diversity.