

# Assignment 5 Complete

March 7, 2022

## Instructions:

- a. You may talk to a friend, discuss the questions and potential directions for solving them. However, you need to write your own solutions and code separately, and not as a group activity.
- b. Do not write your name on the assignment. (1 point)
- c. Export your Jupyter notebook as a PDF file. If you get an error, make sure you have downloaded the MikTeX software (for windows) or MacTeX (for mac). Note that after installing MikTeX/MacTeX, you will need to check for updates, install the updates if needed, and re-start your system. Submit the PDF file. (1 point)
- d. Please include each question (or question number) followed by code and your answer (if applicable). Write your code in the ‘Code’ cells and your answer in the ‘Markdown’ cells of the Jupyter notebook. Ensure that the solution is well-organized, clear, and concise (3 points)
  - 1. It’s easy enough to identify different sections of the homework assignment (e.g., if there are different sections of an assignment, they’re clearly distinguishable by section headers or the like)
  - 2. It’s clear which code/markdown blocks correspond to which questions.
  - 3. There aren’t excessively long outputs of extraneous information (e.g., no printouts of entire data frames without good reason)
- e. This assignment is **due at 11:59pm on Monday, March 7th**. Good luck!
- f. There are two parts in this assignment - Part 1 and Part 2. You have the **option to choose any one part** out of the two and submit the assignment. Each part is worth 35 points. Note that both parts are relevant for the final exam. However, for the purpose of this assignment you need to do any one part of your choice. If you do both parts, only the first part will be graded.

Submissions will be graded with a maximum of **40 points** – 35 points for code & answers, 5 points for anonymity and proper formatting.

## Part 1

The datasets *ENB2012\_Train.csv* and *ENB2012\_Test.csv* provide data on energy analysis using 12 different building shapes simulated in Ecotect. The buildings differ with respect to the glazing area, the glazing area distribution, and the orientation, amongst other parameters.

1. X1 Relative Compactness
2. X2 Surface Area
3. X3 Wall Area
4. X4 Roof Area
5. X5 Overall Height
6. X6 Orientation
7. X7 Glazing Area
8. X8 Glazing Area Distribution
9. y1 Heating Load

For this part, if the 3 criteria suggest different numbers of predictors, consider the best model as per the BIC criterion.

### Q1a

Suppose that we want to implement the best subset selection algorithm to find the first order predictors (X1-X8) that can predict heating load (y1). How many models for  $E(y1)$  are possible, if the model includes (i) one variable, (ii) three variables, and (iii) eight variables? Show your steps without running any code.

Note: The notation  $E(y1)$  means the expected value of y1 or the mean of y1

(3 points for answer)

$$(i) \binom{8}{1} = 8, (ii) \binom{8}{3} = 56, (iii) \binom{8}{8} = 1$$

### Q1b

Implement the best subset selection algorithm to find the “best” first-order predictors of heating load (y1). Print out the model summary.

Use *ENB2012\_Train.csv* and consider only the first-order terms.

(4 points for code)

```
[69]: import pandas as pd
import numpy as np
import statsmodels.formula.api as sm
import seaborn as sns
import matplotlib.pyplot as plt
import itertools
import time
from statsmodels.stats.outliers_influence import variance_inflation_factor
from statsmodels.tools.tools import add_constant
from sklearn.preprocessing import StandardScaler
```

```
from sklearn.linear_model import Ridge, RidgeCV, Lasso, LassoCV
```

```
[3]: ENB_train = pd.read_csv('ENB2012_Train.csv')
      ENB_test = pd.read_csv('ENB2012_Test.csv')
      ENB_train.head()
```

```
[3]:
```

	X1	X2	X3	X4	X5	X6	X7	X8	Y1
0	0.98	514.5	294.0	110.25	7.0	2	0.0	0	15.55
1	0.98	514.5	294.0	110.25	7.0	3	0.0	0	15.55
2	0.98	514.5	294.0	110.25	7.0	5	0.0	0	15.55
3	0.90	563.5	318.5	122.50	7.0	2	0.0	0	20.84
4	0.90	563.5	318.5	122.50	7.0	3	0.0	0	21.46

```
[4]: l= list(ENB_train.columns)
      l.remove('Y1')
      X=ENB_train[l]
```

```
[5]: #Function to develop a model based on all predictors in predictor_subset
      def processSubset(predictor_subset):
          # Fit model on feature_set and calculate R-squared
          model = sm.ols('Y1~' + '+' .join(predictor_subset),data = ENB_train).fit()
          Rsquared = model.rsquared
          return {"model":model, "Rsquared":Rsquared}
```

```
[6]: #Function to select the best model amongst all models with 'k' predictors
      def getBest_model(k):
          tic = time.time()
          results = []
          for combo in itertools.combinations(X.columns, k):
              results.append(processSubset(list(combo)))
          # Wrap everything up in a dataframe
          models = pd.DataFrame(results)
          # Choose the model with the highest RSS
          best_model = models.loc[models['Rsquared'].argmax()]
          toc = time.time()
          print("Processed", models.shape[0], "models on", k, "predictors_
          ↪in", (toc-tic), "seconds.")
          return best_model
```

```
[7]: #Function to select the best model amongst the best models for 'k'_
      ↪predictors, where k = 1,2,3,..
      models_best = pd.DataFrame(columns=["Rsquared", "model"])

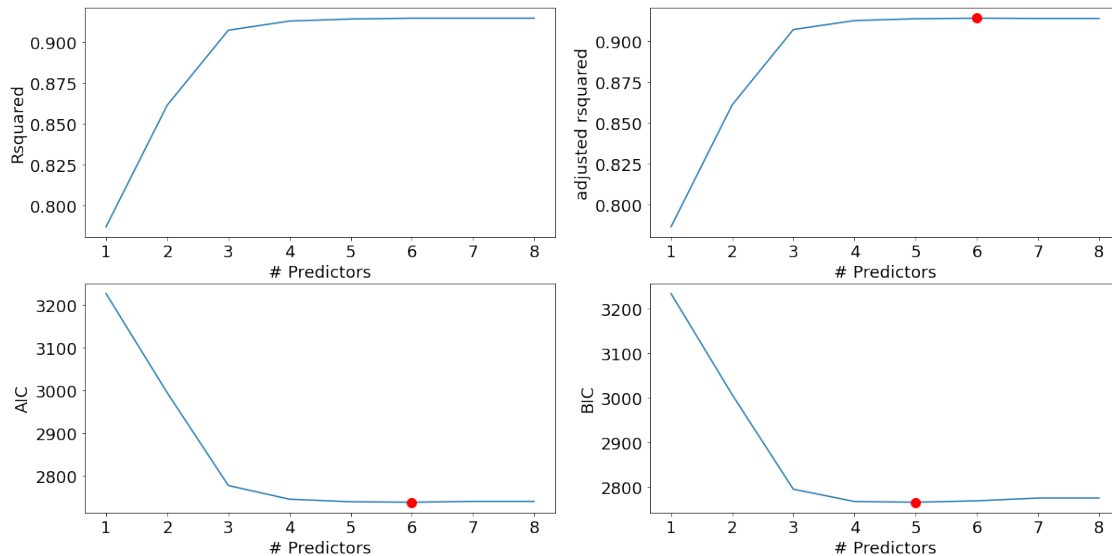
      tic = time.time()
      for i in range(1,1+X.shape[1]):
          models_best.loc[i] = getBest_model(i)
      toc = time.time()
```

```
print("Total elapsed time:", (toc-tic), "seconds.")
```

Processed 8 models on 1 predictors in 0.044007301330566406 seconds.  
Processed 28 models on 2 predictors in 0.11554336547851562 seconds.  
Processed 56 models on 3 predictors in 0.3001894950866699 seconds.  
Processed 70 models on 4 predictors in 0.47556424140930176 seconds.  
Processed 56 models on 5 predictors in 0.43196749687194824 seconds.  
Processed 28 models on 6 predictors in 0.20800161361694336 seconds.  
Processed 8 models on 7 predictors in 0.06399798393249512 seconds.  
Processed 1 models on 8 predictors in 0.00800013542175293 seconds.  
Total elapsed time: 1.686894416809082 seconds.

```
[8]: def best_sub_plots():
    plt.figure(figsize=(20,10))
    plt.rcParams.update({'font.size': 18, 'lines.markersize': 10})
    # Set up a 2x2 grid so we can look at 4 plots at once
    plt.subplot(2, 2, 1)
    # We will now plot a red dot to indicate the model with the largest
    ↪adjusted R2 statistic.
    # The argmax() function can be used to identify the location of the maximum
    ↪point of a vector
    plt.plot(models_best["Rsquared"])
    plt.xlabel('# Predictors')
    plt.ylabel('Rsquared')
    # We will now plot a red dot to indicate the model with the largest
    ↪adjusted R2 statistic.
    # The argmax() function can be used to identify the location of the maximum
    ↪point of a vector
    rsquared_adj = models_best.apply(lambda row: row[1].rsquared_adj, axis=1)
    plt.subplot(2, 2, 2)
    plt.plot(rsquared_adj)
    plt.plot(1+rsquared_adj.argmax(), rsquared_adj.max(), "or")
    plt.xlabel('# Predictors')
    plt.ylabel('adjusted rsquared')
    # We'll do the same for AIC and BIC, this time looking for the models with
    ↪the SMALLEST statistic
    aic = models_best.apply(lambda row: row[1].aic, axis=1)
    plt.subplot(2, 2, 3)
    plt.plot(aic)
    plt.plot(1+aic.argmin(), aic.min(), "or")
    plt.xlabel('# Predictors')
    plt.ylabel('AIC')
    bic = models_best.apply(lambda row: row[1].bic, axis=1)
    plt.subplot(2, 2, 4)
    plt.plot(bic)
    plt.plot(1+bic.argmin(), bic.min(), "or")
    plt.xlabel('# Predictors')
```

```
plt.ylabel('BIC')
best_sub_plots()
```



```
[9]: best_subset_model = models_best['model'][5]
best_subset_model.summary()
```

```
[9]: <class 'statsmodels.iolib.summary.Summary'>
"""
```

```

                                OLS Regression Results
=====
Dep. Variable:                  Y1      R-squared:                0.914
Model:                            OLS      Adj. R-squared:            0.914
Method:                 Least Squares      F-statistic:                1144.
Date:                Thu, 24 Feb 2022      Prob (F-statistic):        2.87e-283
Time:                  03:39:33      Log-Likelihood:            -1363.5
No. Observations:                542      AIC:                       2739.
Df Residuals:                    536      BIC:                       2765.
Df Model:                          5
Covariance Type:                nonrobust
=====

```

	coef	std err	t	P> t	[0.025	0.975]
Intercept	46.1893	5.212	8.862	0.000	35.950	56.428
X1	-45.8814	2.567	-17.876	0.000	-50.923	-40.839
X4	-0.1078	0.014	-7.539	0.000	-0.136	-0.080
X5	4.7438	0.318	14.902	0.000	4.118	5.369
X7	20.1466	0.990	20.357	0.000	18.203	22.091
X8	0.2421	0.086	2.809	0.005	0.073	0.411

```
=====
Omnibus:                9.613    Durbin-Watson:                0.869
Prob(Omnibus):          0.008    Jarque-Bera (JB):          15.883
Skew:                  -0.031    Prob(JB):                  0.000356
Kurtosis:              3.836    Cond. No.                  7.80e+03
=====
```

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

[2] The condition number is large, 7.8e+03. This might indicate that there are strong multicollinearity or other numerical problems.

"""

The best predictors according to best subset selection algorithm are x1, x2, x3, x5, x7, x8.

## Q1c

Should R-squared be used to select from among a set of models with different numbers of predictors? Justify your answer.

*(1 point for answer, 2 points for justification)*

No, R-squared cannot be used to select from among a set of models with different numbers of predictors. Since R-squared increases monotonically as the number of predictors increase, so it would always lead to selecting the model with all the predictors.

## Q1d

Calculate the RMSE of the model found in (b). Compare it with the RMSE of the model using all first-order predictors. You will find that the two RMSEs are similar. Seems like the best subset model didn't help improve prediction. However, did it help improve inference? Justify your answer.

Hint: VIF!

*(2 points for code, 2 points for answer)*

```
[10]: pred_y1 = best_subset_model.predict(ENB_test)
      np.sqrt(((pred_y1 - ENB_test.Y1)**2).mean())
```

```
[10]: 2.778167195091445
```

```
[11]: model = sm.ols('Y1~' + '+' .join(X.columns), data = ENB_train).fit()
      pred_y1 = model.predict(ENB_test)
      np.sqrt(((pred_y1 - ENB_test.Y1)**2).mean())
```

```
[11]: 2.774159872401328
```

```
[12]: model.summary()
```

```
[12]: <class 'statsmodels.iolib.summary.Summary'>
```

```
"""
                                OLS Regression Results
=====
Dep. Variable:                  Y1      R-squared:                0.915
Model:                        OLS      Adj. R-squared:            0.914
Method:                    Least Squares   F-statistic:                819.5
Date:                Thu, 24 Feb 2022   Prob (F-statistic):        6.84e-281
Time:                        03:40:46   Log-Likelihood:            -1361.9
No. Observations:                542     AIC:                       2740.
Df Residuals:                    534     BIC:                       2774.
Df Model:                          7
Covariance Type:                nonrobust
=====
              coef      std err          t      P>|t|      [0.025      0.975]
-----
Intercept      86.8999      23.677        3.670      0.000      40.389     133.410
X1             -68.0136      12.782       -5.321      0.000     -93.122     -42.905
X2              -0.0642       0.017       -3.842      0.000       -0.097      -0.031
X3               0.0363       0.005        7.617      0.000        0.027      0.046
X4             -0.0503       0.009       -5.335      0.000       -0.069     -0.032
X5               4.2704       0.417       10.243      0.000        3.451      5.089
X6               0.0223       0.118        0.189      0.850       -0.209      0.254
X7             20.1554       0.989       20.386      0.000       18.213     22.098
X8               0.2430       0.086        2.822      0.005        0.074      0.412
=====
Omnibus:                 10.570   Durbin-Watson:           0.892
Prob(Omnibus):             0.005   Jarque-Bera (JB):        18.131
Skew:                     -0.039   Prob(JB):                 0.000116
Kurtosis:                  3.893   Cond. No.                 2.31e+16
=====
```

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

[2] The smallest eigenvalue is 5.97e-25. This might indicate that there are strong multicollinearity problems or that the design matrix is singular.

```
"""
```

The RMSE of the best subset model is close to the RMSE of the model with all the predictors. Note that X2, X3 and X6 are excluded in the best subset model. X6 is statistically insignificant, so it is not a useful predictor and is removed in the best subset. If we compute the VIF with either X2 or X3 included in the model, we find that X2 and X3 are multicollinear with other predictors in the model. Thus, they do not help explain any additional variation in the response, given other predictors. So, they are not necessary and that is why not there in the best subset model. Thus, the best subset model eliminates the issue of multicollinearity, and provides a more accurate inference.

```
[13]: def vif(X_vif):
        X_vif = add_constant(X_vif)
        vif_data = pd.DataFrame()
        vif_data["feature"] = X_vif.columns

        for i in range(len(X_vif.columns)):
            vif_data.loc[i, 'VIF'] = variance_inflation_factor(X_vif.values, i)

        print(vif_data)
```

```
[14]: #VIF with X2 in the model
X_vif=X.iloc[:,[0,1,3,4,6,7]]
vif(X_vif)
```

	feature	VIF
0	const	33578.814492
1	X1	110.732796
2	X2	116.487267
3	X4	33.819552
4	X5	31.863278
5	X7	1.050550
6	X8	1.048236

```
[15]: #VIF with X3 in the model
X_vif=X.iloc[:,[0,2,3,4,6,7]]
vif(X_vif)
```

	feature	VIF
0	const	33578.814492
1	X1	110.732796
2	X3	26.876039
3	X4	220.820817
4	X5	31.863278
5	X7	1.050550
6	X8	1.048236

```
[16]: #VIF with X2 and X3 removed from the model
X_vif=X.iloc[:,[0,3,4,6,7]]
vif(X_vif)
```

	feature	VIF
0	const	1623.968159
1	X1	4.457484
2	X4	24.792621
3	X5	18.546746
4	X7	1.050511
5	X8	1.048155



## Q1e

Let us consider adding all the 2-factor interactions of the predictors in the model. Please answer the following questions without running code.

- i) How many predictors do we have in total?
- ii) Assume best subset selection is used. How many models are fitted in total?
- iii) Assume forward selection is used. How many models are fitted in total?
- iv) Assume backward selection is used. How many models are fitted in total?
- v) How many models will be developed in the iteration that contains exactly 10 predictors in each model – for best subsets, fwd/bwd regression?
- vi) What approach would you choose for variable selection (amongst best subset, forward selection, backward selection)?

(8 points for answer)

i)  $8 + \binom{8}{2} = 8 + 28 = 36$

ii)  $2^p = 2^{36} = 68,719,476,736$

iii)  $1 + p(p+1)/2 = 1 + 36 * 37/2 = 667$

iv)  $1 + p(p+1)/2 = 1 + 36 * 37/2 = 667$

v) Best subset: 10 predictors,  $\binom{36}{10} = 254,186,856$

Fwd selection: 10 predictors, p-k models for k+1 predictors, then 10 predictors corresponds to  $k = 9$ ,  $p - k = 36 - 9 = 27$

Bwd selection: 10 predictors, we consider k models that contains k+1 predictors, we have  $10 + 1 = 11$  models that contains 10 predictors

- vi) We should use forward or backward selection because best subset selection is very computationally heavy in this case.

## Q1f

Use forward selection to find the “best” first-order predictors and 2-factor interactions of the predictors of y1 (Heating Load). Print out the model summary.

(5 points for code)

```
[20]: X_backup=ENB_train[1]

for combo in itertools.combinations(X_backup.columns, 2):
    ENB_train['_'.join(combo)] = ENB_train[combo[0]]*ENB_train[combo[1]]
    ENB_test['_'.join(combo)] = ENB_test[combo[0]]*ENB_test[combo[1]]
    X.loc[:, '_'.join(combo)] = ENB_train.loc[:, '_'.join(combo)]

[19]: #Function to find the best predictor out of p-k predictors and add it to the
      ↪ model containing the k predictors
def forward(predictors):
```

```

# Pull out predictors we still need to process
remaining_predictors = [p for p in X.columns if p not in predictors]
tic = time.time()
results = []
for p in remaining_predictors:
    results.append(processSubset(predictors+[p]))
# Wrap everything up in a nice dataframe
models = pd.DataFrame(results)
# Choose the model with the highest RSS
best_model = models.loc[models['Rsquared'].argmax()]
toc = time.time()
print("Processed ", models.shape[0], "models on",
len(predictors)+1, "predictors in", (toc-tic), "seconds.")
# Return the best model, along with some other useful information about the
model
return best_model

```

```

[21]: def forward_selection():
    models_best = pd.DataFrame(columns=["Rsquared", "model"])
    tic = time.time()
    predictors = []
    for i in range(1, len(X.columns)+1):
        models_best.loc[i] = forward(predictors)
        predictors = list(models_best.loc[i]["model"].params.index[1:])
    toc = time.time()
    print("Total elapsed time:", (toc-tic), "seconds.")
    return models_best

```

```

[22]: models_best = forward_selection()

```

```

Processed 36 models on 1 predictors in 0.197035551071167 seconds.
Processed 35 models on 2 predictors in 0.15621089935302734 seconds.
Processed 34 models on 3 predictors in 0.1815474033355713 seconds.
Processed 33 models on 4 predictors in 0.2255539894104004 seconds.
Processed 32 models on 5 predictors in 0.2816908359527588 seconds.
Processed 31 models on 6 predictors in 0.2708756923675537 seconds.
Processed 30 models on 7 predictors in 0.2542533874511719 seconds.
Processed 29 models on 8 predictors in 0.328472375869751 seconds.
Processed 28 models on 9 predictors in 0.3839991092681885 seconds.
Processed 27 models on 10 predictors in 0.31999826431274414 seconds.
Processed 26 models on 11 predictors in 0.3524012565612793 seconds.
Processed 25 models on 12 predictors in 0.3279991149902344 seconds.
Processed 24 models on 13 predictors in 0.3922576904296875 seconds.
Processed 23 models on 14 predictors in 0.316530704498291 seconds.
Processed 22 models on 15 predictors in 0.34543752670288086 seconds.
Processed 21 models on 16 predictors in 0.3589646816253662 seconds.
Processed 20 models on 17 predictors in 0.37851786613464355 seconds.
Processed 19 models on 18 predictors in 0.4036996364593506 seconds.

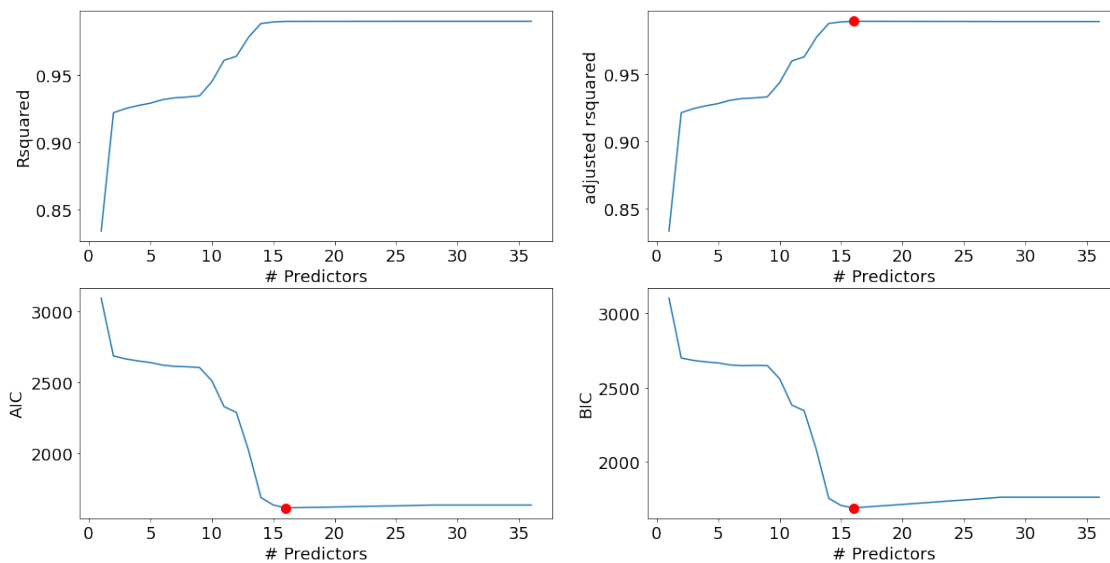
```

```

Processed 18 models on 19 predictors in 0.40822577476501465 seconds.
Processed 17 models on 20 predictors in 0.38858795166015625 seconds.
Processed 16 models on 21 predictors in 0.37599873542785645 seconds.
Processed 15 models on 22 predictors in 0.39428043365478516 seconds.
Processed 14 models on 23 predictors in 0.34958434104919434 seconds.
Processed 13 models on 24 predictors in 0.32555556297302246 seconds.
Processed 12 models on 25 predictors in 0.3119509220123291 seconds.
Processed 11 models on 26 predictors in 0.2833878993988037 seconds.
Processed 10 models on 27 predictors in 0.3082747459411621 seconds.
Processed 9 models on 28 predictors in 0.24689698219299316 seconds.
Processed 8 models on 29 predictors in 0.2558615207672119 seconds.
Processed 7 models on 30 predictors in 0.2569773197174072 seconds.
Processed 6 models on 31 predictors in 0.20799875259399414 seconds.
Processed 5 models on 32 predictors in 0.1919999122619629 seconds.
Processed 4 models on 33 predictors in 0.11999821662902832 seconds.
Processed 3 models on 34 predictors in 0.11210274696350098 seconds.
Processed 2 models on 35 predictors in 0.07199883460998535 seconds.
Processed 1 models on 36 predictors in 0.03196454048156738 seconds.
Total elapsed time: 10.258779764175415 seconds.

```

```
[23]: best_sub_plots()
```



```

[24]: #all criteria agree on having 16 predictors
best_fwd_reg_model = models_best['model'][16]
best_fwd_reg_model.summary()

```

```

[24]: <class 'statsmodels.iolib.summary.Summary'>
      """

```

OLS Regression Results

```

=====
Dep. Variable:          Y1      R-squared:          0.990
Model:                  OLS      Adj. R-squared:       0.989
Method:                 Least Squares      F-statistic:       3118.
Date:                   Thu, 24 Feb 2022      Prob (F-statistic):    0.00
Time:                   03:42:56      Log-Likelihood:       -792.42
No. Observations:       542      AIC:                 1619.
Df Residuals:           525      BIC:                 1692.
Df Model:               16
Covariance Type:        nonrobust
=====

```

```

=====
              coef      std err          t      P>|t|      [0.025      0.975]
-----
Intercept    -7705.1653     607.596    -12.681     0.000    -8898.783    -6511.548
X3_X5         1.5109       0.109     13.891     0.000         1.297         1.725
X5_X7        -2.7764       0.830     -3.346     0.001        -4.407        -1.146
X1_X2        -4.9300       0.652     -7.560     0.000        -6.211        -3.649
X2_X5        -0.9502       0.067    -14.222     0.000        -1.081        -0.819
X3_X8         0.0003       0.000         0.671     0.502        -0.001         0.001
X7_X8        -2.8188       0.204    -13.814     0.000        -3.220        -2.418
X7           84.3219       9.950      8.475     0.000        64.775       103.868
X4_X7        -0.2434       0.032     -7.592     0.000        -0.306        -0.180
X3_X4         0.0423       0.002     17.156     0.000         0.037         0.047
X3           7.2024       1.546      4.660     0.000         4.166        10.239
X1_X5        196.7862      36.713      5.360     0.000       124.663       268.910
X4           55.3400       3.229     17.139     0.000         48.997         61.683
X1          2968.1523     606.145      4.897     0.000      1777.385      4158.919
X2_X4        -0.0460       0.004    -12.152     0.000        -0.053        -0.039
X2_X3        -0.0047       0.001     -7.500     0.000        -0.006        -0.003
X1_X8         0.9009       0.191      4.723     0.000         0.526         1.276
=====

```

```

=====
Omnibus:          1.817      Durbin-Watson:          0.832
Prob(Omnibus):    0.403      Jarque-Bera (JB):        1.620
Skew:             0.095      Prob(JB):                0.445
Kurtosis:         3.188      Cond. No.                4.84e+09
=====

```

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

[2] The condition number is large, 4.84e+09. This might indicate that there are strong multicollinearity or other numerical problems.

"""

## Q1g

Calculate the RMSE of the model found in (f). Compare it with i) the RMSE of model you found in (b) and ii) the RMSE of the model using all first-order and 2-factor interaction terms and discuss about your finding.

*(2 points for code, 2 points for answer)*

```
[51]: full_model = models_best['model'][36]
```

```
[26]: pred_y1 = best_fwd_reg_model.predict(ENB_test)
      np.sqrt(((pred_y1 - ENB_test.Y1)**2).mean())
```

```
[26]: 1.0894404698394693
```

The RMSE is much smaller than the RMSE of the model we found in (b), which is 2.7781671950919242. We can see a huge improvement of both the adjusted R-squared (0.914 to 0.989) and RMSE (1.09 to 2.78) after adding the 2-factor interaction terms.

```
[27]: #calculate RMSE for all predictors
      pred_y1 = full_model.predict(ENB_test)
      np.sqrt(((pred_y1 - ENB_test.Y1)**2).mean())
```

```
[27]: 1.1088320527269544
```

We can see that the best model RMSE is very close to the RMSE of the full model. However, there are several insignificant predictors in the full model, which may be due to multicollinearity, or just due to the fact that those predictors do not explain the response. Thus, the model in (f) is better for the purpose of inference.

## Q1h

Assume that we found another dataset of 32 variables on the same set of 768 buildings (542 for training) that we would want to add into our model. We want find the “best” model of all 40 predictors and their 2-factor interaction terms. Would you choose forward or backward selection? Justify your answer.

*(1 point for answer, 3 points for justification)*

$p = 40 + \binom{40}{2} = 40 + 780 = 820$  which is larger than  $n = 542$ . We should use forward selection in this case since backward selection requires that the number of samples  $n$  is larger than the number of variables  $p$  (so that the full model can be fit). In contrast, forward stepwise can be used even when  $n < p$ , and so is the only viable subset method when  $p$  is very large.

## Part 2

See <https://exoplanetarchive.ipac.caltech.edu> (for context/source). We are using the Composite Planetary Systems dataset

## Q2a

Say we're interested in modeling the radius of exoplanets in kilometers, which is named as `pl_rade` in the data. Note that the variable `pl_rade` captures the radius of each plant as a proportion of Earth's radius, which is approximately 6,378.1370 km.

Develop a linear regression model to predict `pl_rade` using all the variables in `train_CompositePlanetarySystems.csv` except `pl_name`, `disc_facility` and `disc_locale`. Find the RMSE (Root mean squared error) of the model on `test1_CompositePlanetarySystems.csv` and `test2_CompositePlanetarySystems.csv`.

(4 points for code)

[82]: train

```
[82]:      sy_snum  sy_pnum  cb_flag  disc_year  rv_flag  pul_flag  ptv_flag  \
0          1         2         0        2014         0         0         0
1          1         3         0        2019         0         0         0
2          1         1         0        2021         0         0         0
3          1         1         0        2020         0         0         0
4          1         5         0        2014         0         0         0
..      ...      ...      ...      ...      ...      ...      ...
965         1         1         0        2021         0         0         0
966         1         2         0        2013         0         0         0
967         1         3         0        2012         0         0         0
968         1         2         0        2016         0         0         0
969         1         1         0        2021         0         0         0

      tran_flag  ast_flag  obm_flag  ...  ima_flag  pl_controv_flag  \
0          1         0         0  ...         0         0
1          1         0         0  ...         0         0
2          1         0         0  ...         0         0
3          1         0         0  ...         0         0
4          1         0         0  ...         0         0
..      ...      ...      ...  ...      ...      ...
965         1         0         0  ...         0         0
966         1         0         0  ...         0         0
967         1         0         0  ...         0         0
968         1         0         0  ...         0         0
969         1         0         0  ...         0         0

      pl_orbper  pl_dens  ttv_flag  st_teff  st_rad  st_mass  st_lum  st_logg
0    10.613839    3.040         0    5818.0    0.89    0.91  -0.331    4.54
1     3.359605    4.980         0    5212.9    0.76    0.80  -0.401    4.60
2     6.892130    5.600         0    6164.0    1.45    1.29   0.437    4.23
3     3.262000    2.480         0    3752.0    0.51    0.52  -1.333    4.74
4    87.090195    2.300         0    4997.0    0.76    0.84  -0.496    4.58
..      ...      ...      ...      ...      ...      ...      ...
965    4.626050    2.970         0    5834.0    0.80    0.85  -0.075    4.57
```

966	191.231800	0.152	1	5600.0	1.82	1.10	0.326	3.96
967	85.312000	0.030	1	6018.0	0.94	1.04	-0.142	4.51
968	8.752900	4.310	0	4326.0	0.55	0.61	-0.874	4.74
969	5.186720	2.420	0	5230.0	0.78	0.87	-0.394	4.59

[970 rows x 22 columns]

```
[94]: train = pd.read_csv("train_CompositePlanetarySystems.csv")
test1 = pd.read_csv("test1_CompositePlanetarySystems.csv")
test2 = pd.read_csv("test2_CompositePlanetarySystems.csv")
```

```
[95]: def data_prep(data):
        X = data.drop(columns = [
            'pl_name', 'pl_rade', 'disc_facility', 'disc_locale'], axis=1)
        return X
X = data_prep(train)
Xtest1 = data_prep(test1)
Xtest2 = data_prep(test2)
```

```
[96]: model = sm.ols('pl_rade~' + '+' .join(X.columns), data = train).fit()
```

```
[97]: def rmse(data, ytrue, model):
        pred = model.predict(data)
        rmse = np.sqrt(((pred-ytrue)**2).mean())
        return rmse
```

```
[98]: rmse(test1, test1.pl_rade, model)
```

```
[98]: 127.36595496362311
```

RMSE on *test1.csv* is 127

```
[99]: rmse(test2, test2.pl_rade, model)
```

```
[99]: 202.66277572552076
```

RMSE on *test2.csv* is 203

## Q2b

Develop a ridge regression model to predict `pl_rade` using all the variables in *train\_CompositePlanetarySystems.csv* except `pl_name`, `disc_facility` and `disc_locale`. What is the optimal value of the tuning parameter  $\lambda$ ?

**Hint:** You may use the following grid of lambda values to find the optimal  $\lambda$ : `alphas = 10**np.linspace(2, 0.5, 200)*0.5`

Remember to standardize data before fitting the ridge regression model

(5 points for code)

```
[36]: #Standardizing predictors so that each of them have zero mean and unit variance

#Defining a scaler object
scaler = StandardScaler()

#The scaler object will contain the mean and variance of each column
↳ (predictor) of X. These values will be useful to scale
#test data based on the same mean and variance as obtained on train data
scaler.fit(X)

#Using the scaler object (or the values of mean and variance stored in it) to
↳ standardize X (or train data) and other datasets
Xstd = pd.DataFrame(scaler.transform(X))
Xtest1_std = pd.DataFrame(scaler.transform(Xtest1))
Xtest2_std = pd.DataFrame(scaler.transform(Xtest2))
Xstd.columns = X.columns
Xtest1_std.columns = Xtest1.columns
Xtest2_std.columns = Xtest2.columns

y = train.pl_rade
```

```
[60]: #Let us use cross validation to find the optimal value of the tuning parameter
↳ - lambda
#For the optimal lambda, the cross validation error will be the least

alphas = 10*np.linspace(2,0.5,200)*0.5
ridgecv = RidgeCV(alphas = alphas,store_cv_values=True)
ridgecv.fit(Xstd, y)

#Optimal value of the tuning parameter - lamda
ridgecv.alpha_
```

[60]: 3.221181754360687

The optimal value of  $\lambda$  is 3.22

## Q2c

Use the optimal value of  $\lambda$  found in the previous question to develop a ridge regression model. What is the RMSE of the model on *test1\_CompositePlanetarySystems.csv* and *test2\_CompositePlanetarySystems.csv*?

(5 points for code)

```
[61]: #Using the developed ridge regression model to predict on test data
ridge = Ridge(alpha = ridgecv.alpha_)
ridge.fit(Xstd, y)
```



```
[61]: Ridge(alpha=3.221181754360687)
```

```
[62]: rmse(Xtest1_std,test1.pl_rade,ridge)
```

```
[62]: 3.218708879148552
```

```
[63]: rmse(Xtest2_std,test2.pl_rade,ridge)
```

```
[63]: 2.9898651678534853
```

The RMSE on *test1.csv* is 3.22 and on *test2.csv* is 2.99.

## Q2d

Note that ridge regression has a much lower RMSE on test datasets as compared to Ordinary least squares (OLS) regression. Shrinking the coefficients has improved the model fit. Appreciate it. Which are the top two predictors for which the coefficients have shrunk the most?

To answer this question, find the ridge regression estimates for  $\lambda = 10^{-10}$ . Treat these estimates as OLS estimates and find the predictors for which these estimates have shrunk the most in the model developed in 2c.

*(4 points for code, 1 point for answer)*

```
[41]: ridge_ols = Ridge(alpha = 10**-10)
      ridge_ols.fit(Xstd, y)
```

```
[41]: Ridge(alpha=1e-10)
```

```
[42]: coef_diff = pd.Series(np.abs(ridge_ols.coef_) - np.abs(ridge.coef_))
```

```
[43]: pd.concat([pd.Series(Xstd.columns),coef_diff],axis=1).sort_values(by =
      ↪1,ascending=False)[0:2]
```

```
[43]:           0           1
      14  pl_orbper  75.809870
      12   ima_flag  75.800798
```

The coefficients of *pl\_orbper* and *ima\_flag* have shrunk the most.

## Q2e

Why do you think the coefficients of the two variables identified in the previous question shrunk the most?

**Hint:** VIF!

*(4 points for justification - including any code used)*

```
[44]: def vif(X):
      X = add_constant(X)
      vif_data = pd.DataFrame()
      vif_data["feature"] = X.columns

      for i in range(len(X.columns)):
          vif_data.loc[i, 'VIF'] = variance_inflation_factor(X.values, i)

      print(vif_data)
```

```
[45]: vif(X)
```

	feature	VIF
0	const	293592.094382
1	sy_snum	1.162755
2	sy_pnum	1.146094
3	cb_flag	2.058390
4	disc_year	1.362591
5	rv_flag	2.014735
6	pul_flag	NaN
7	ptv_flag	2.610996
8	tran_flag	2.466560
9	ast_flag	1.091332
10	obm_flag	1.183595
11	micro_flag	NaN
12	etv_flag	7.440160
13	ima_flag	37675.155715
14	pl_controv_flag	1.698968
15	pl_orbper	37683.517115
16	pl_dens	1.017593
17	ttv_flag	1.110245
18	st_teff	13.944718
19	st_rad	3.257504
20	st_mass	4.139998
21	st_lum	12.507339
22	st_logg	10.421923

```
C:\Users\akl0407\Anaconda3\lib\site-
packages\statsmodels\regression\linear_model.py:1715: RuntimeWarning: invalid
value encountered in double_scalars
  return 1 - self.ssr/self.centered_tss
```

Note that `ima_flag` and `pl_orbper` are highly multicollinear, which results in a high variance of their coefficients. Thus, their coefficients have been shrunk the most in the ridge regression model, as shrinkage reduces the variance of the coefficients to improve the model fit.

## Q2f

Develop a lasso model to predict `pl_rade` using all the variables in `train_CompositePlanetarySystems.csv` except `pl_name`, `disc_facility` and `disc_locale`. What is the optimal value of the tuning parameter  $\lambda$ ?

**Hint:** You may use the following grid of lambda values to find the optimal  $\lambda$ : `alphas = 10**np.linspace(0,-2.5,200)*0.5`

(4 points for code)

```
[64]: #Let us use cross validation to find the optimal value of the tuning parameter
      ↪- lambda
      #For the optimal lambda, the cross validation error will be the least

      alphas = 10**np.linspace(0,-2.5,200)*0.5
      lassocv = LassoCV(alphas = alphas, cv = 10, max_iter = 100000)
      lassocv.fit(Xstd, y)

      #Optimal value of the tuning parameter - lamda
      lassocv.alpha_
```

```
[64]: 0.002051329052913595
```

Optimal  $\lambda$  for lasso is 0.002

## Q2g

Use the optimal value of  $\lambda$  found in the previous question to develop a lasso model. What is the RMSE of the model on `test1_CompositePlanetarySystems.csv` and `test2_CompositePlanetarySystems.csv`?

(5 points for code)

```
[65]: #Using the developed lasso model to predict on test data
      lasso = Lasso(alpha = lassocv.alpha_)
      lasso.fit(Xstd, y)
```

```
[65]: Lasso(alpha=0.002051329052913595)
```

```
[66]: rmse(Xtest1_std,test1.pl_rade,lasso)
```

```
[66]: 3.220582381965326
```

```
[67]: rmse(Xtest2_std,test2.pl_rade,lasso)
```

```
[67]: 2.996713163709016
```

## Q2h

Note that lasso has a much lower RMSE on test datasets as compared to Ordinary least squares (OLS) regression. Shrinking the coefficients has improved the model fit. Appreciate it. Which variables have been eliminated by lasso?

To answer this question, find the predictors whose coefficients are 0 in the lasso model.

*(2 points for code, 1 point for answer)*

```
[68]: Xstd.columns[lasso.coef_==0]
```

```
[68]: Index(['pul_flag', 'micro_flag', 'ima_flag'], dtype='object')
```

Coefficients of `pul_flag`, `micro_flag`, and `ima_flag` are reduced to zero. Thus, these predictors are eliminated by lasso.